The Impact of Deep Learning on Speech Synthesis with Mobile Devices

Hannes Bohnengel

Technical University of Munich

hannes.bohnengel@tum.de

ABSTRACT

Statistical Parametric Speech Synthesis (SPSS) is a commonly used approach to synthesize speech and has several advantages over the alternative methods like formant-based or unit-selection synthesis. The two most significant advantages are the small model size and the easily adjustable voice characteristics. However, the quality of the generated speech still has potential for improvement. Due to over-smoothing the voice sounds muffled in comparison to natural speech [5]. To address this issue, the use of deep learning models has been suggested by different researchers [11, 17]. Throughout their experiments, a more natural voice as well as an improved prediction performance could be achieved. When implementing speech synthesis on mobile devices, there arise different issues like memory constraints of mobile operating systems or the claim for real-time responsiveness. To tackle these challenges, two approaches have been introduced. The first approach includes a number of optimization steps and results in a decrease of the computation time by 65 % without a considerable loss of voice quality [15]. In the second approach a deep learning model has been deployed in different parts of a SPSS system [6]. Here a tremendous reduction of the footprint size is the main result and that way enables the independence on Internet-based applications. In general, it can be concluded that the use of deep learning models for speech synthesis on mobile devices is a reasonable measure.

1 INTRODUCTION

Virtual Personal Assistants (VPAs) like Siri, Cortana or Google Now start having a huge impact on the way of interacting with electronic devices like smartphones or notebooks. Up to now the VPAs help with rather simple tasks such as search queries, starting phone calls or setting a clock, but according to a recent survey from the IT research firm, Gartner [4], this will change in the near future. With the Facebook Messenger it is already possible to make purchases or to order an Uber car and new use cases are expected soon. The survey also states that through the vast increase of devices in the scope of the Internet of Things (IoT) the way of interacting with machines will go towards minimal or zero touch. Instead of interacting through common touch-displays or buttons, the user simply speaks to the device like to another person. To enable this, both speech recognition and speech synthesis are essential technologies.

In this paper, I will only focus on the speech synthesis part. A widely spread technique to synthesize human speech from a given text or from linguistic descriptions is Statistical Parametric Speech Synthesis (SPSS), also referred to as Statistical Parametric Speech Generation (SPSG) [13]. This technique is based on the usage of Hidden Markov Models (HMMs). Black *et al.* [5] have shown that it has several advantages over its predecessor, the concatenative speech synthesis, for example the flexibility in changing voice characteristics and a smaller memory footprint. However, the quality of the generated speech still has potential for improvement. Due to over-smoothing the voice sounds muffled in comparison to natural speech.

This is where recent achievements in deep learning come in. Deep learning is usually referred to as a class of machine learning tech-niques that achieve tasks like feature extraction or pattern analysis by using many connected layers of non-linear information processing [9, 13]. Since 2006, advances in the training algorithms of Deep Neural Networks (DNNs) have enabled the field of deep learning applications to emerge [6]. Most machine learning models until then had used shallow structures, like HMMs, Gaussian Mixture Models (GMMs), Conditional Random Fields (CRFs) or Support Vector Machines (SVMs). In these structures, only one layer is responsible for generating features out of the raw input signals. While achieving quite good results with rather simple problems, they reach their limit when it comes to more complex tasks like processing human language or natural images [9]. In the tutorial survey [9], the author also states four different approaches to improve speech synthesis through deep learning models, whereof three are dealing with SPSS. One of those three approaches is described in [17], where the authors implemented a part of the speech synthesis system by using a DNN and observed an improved performance in predicting output features. In [11], a more general approach is conducted by investigating what effects the deployment of a DNN on different parts of the SPSS system has. An improvement of the naturalness of the generated speech was one of the main results.

For implementing speech synthesis on resource-constrained devices like smartphones or tablets, SPSS is considered the best solution due to the trade-off between voice quality and acceptable footprint size [15]. Since the computational costs of SPSS are often high, some optimization steps are conducted in [15] to make the HMM-based speech synthesis technique SPSS more suitable for mobile devices. These steps include reducing the size of the decision trees and introducing streaming synthesis. Going one step further, in [6] an approach to adapt SPSS for embedded devices by using a deep learning model, an auto encoder, is employed. Four tasks (syllabification, phonetic transcription, part-of-speech tagging and lexical stress prediction) are examined and tested with the use of this deep learning model. As results, the authors highlight hugely reduced model sizes and a very close performance to the state-of-the-art models. This shows, that the usage of deep learning models for speech synthesis on embedded systems is a reasonable step, not only to improve performance and voice quality, but also towards the independence of online databases for speech synthesis.

The remaining paper is structured as follows: Section 2 first states the motivation by explaining why speech synthesis is a useful technology. Then it describes the conventional approach without deep learning models for speech synthesis and gives an overview of advantages and drawbacks of the used models and techniques. This is followed by a brief explanation of one of the most commonly used technique, SPSS, where the paper [5] has been chosen as typically cited reference. Thereafter, two possibilities how SPSS can be improved by deploying deep learning models are characterized. For this purpose, the papers [17] and [11] are reviewed. In Section 4, the motivation why speech synthesis is important on mobile devices is given, followed by two examples on how speech synthesis can be implemented on an embedded system, one without [15] and one with deep learning models [6]. Finally, Section 5 summarizes the essential points of this paper and gives some future directions.

2 CONVENTIONAL SPEECH SYNTHESIS

2.1 Motivation and Approaches

Speech synthesis has emerged over the last ten years due to a vast contribution by the global community of researchers and the increasing computational power for data processing. Its quality and naturalness has increased steadily and different approaches have been developed so far [14]. The typical applications like navigation systems in cars or telephone-based dialogue systems are nowadays widely established. But also as reading aid for visually impaired people [1] or as in the case of the famous scientist Stephen Hawking, who has been using a synthesized voice to communicate since 1997 [3], speech synthesis has proven to be very useful. Another very interesting application of speech synthesis is shown in [10]. Here, the author proposed to introduce synthetic speech as means of communication between pilots, since there have been many accidents due to misunderstandings at radio-based communication.

According to [12], speech synthesis can be divided into three types: Canned speech, Context-to-Speech (CTS) and Text-to-Speech (TTS). Canned speech more or less is the playback of prerecorded spoken sentences or words, both with none or very little adjustments. Typical examples are the announcements on train stations. Because of the high effort of recording everything (almost) exactly as it is played back, this approach is limited to a few simple applications. With the second type, CTS, the waveform is generated out of a linguistic description without any information of the respective text. In this way, no natural language processing is required, but nevertheless CTS has not made any important impact yet. The last and most promising type is TTS. A TTS system consists of a Natural Language Processing (NLP) part, where the text is analyzed and the word and sentence structure and accents are extracted. In the next step, these accents are used to generate the prosody of the given text like duration, intensity and pitch. Then, the created phonetic representations with prosody information are stringed together to a continuous stream of signal parameters. The last task, the speech generation, uses this stream to generate the respective waveform. This function block can be implemented in different ways. In [12], three general approaches are named as follows: Parametric Speech Synthesis (formant-based synthesis), Concatenative Speech Synthesis (unit-selection synthesis) and Statistical Parametric Speech Synthesis (SPSS) (Hidden Markov Model (HMM)-based synthesis). The methods in brackets are the respective implementations, which are most commonly used.

The formant-based synthesis is the oldest approach. To generate a voice waveform, an excitation signal is fed into multiple formant filters, which describe the characteristics of the human vocal tract. The output of the filters then forms the voice waveform. This technique is the only one which does not need any recorded speech, but instead generates the synthesized voice only by modeling the human vocal tract. On the one hand, the quality of the generated voice is the lowest in comparison to the other techniques, on the other hand, formant-based synthesizers have the smallest footprint and their voice characteristics can easily been modified by just changing their filter parameters [12].

With the development of concatenative speech synthesizer, the quality of the generated speech improved tremendously. Very similar to CTS prerecorded speech is used as reference. Basically, the recorded speech is divided into units and these units are then stringed together to form the new speech signal according to a given text. Hereby, the chosen size of the units determines both the footprint size and the voice quality. With larger units a higher voice quality can be achieved, but this also results in a much larger database. The challenge in unit-selection is to ensure that the transitions between the units are as natural as possible [12]. In Section 2.2, the HMM-based synthesis, a specific instance of SPSS is described in detail.

2.2 HMM-based Speech Synthesis

In this section, the most recent approach, *i.e.* HMM-based speech synthesis, will be described further and both advantages and drawbacks compared to unit-selection synthesis will be highlighted. Therefore, the work of Black *et al.* [5] will be taken as reference.

The quality of unit-selection synthesis directly depends on the quality of the prerecorded speech. But even with a database of excellent quality, sporadic errors can still not be avoided totally. If a specific phonetic or prosodic part of a generated sentence is not well represented in the database the output quality of this sentence suffers immensely. To try to avoid this, a huge effort in specifically designing the database for the required application can be performed, but still there is no guarantee that such bad joins will not happen. In addition, the fact that in unit-selection no or only very little adaptions of the voice characteristics are possible without an enormous increase of the database size, the ambition towards seamless speech synthesis leads towards the HMM-based approach. Here, a statistic representation of some sets of speech segments is used to generate arbitrary synthetic speech.

In Figure 1, the structure of a typical HMM-based synthesizer is shown. The whole system can be divided into two parts, the training and the synthesis part. The connection of these two parts is a set of context-dependent HMMs.
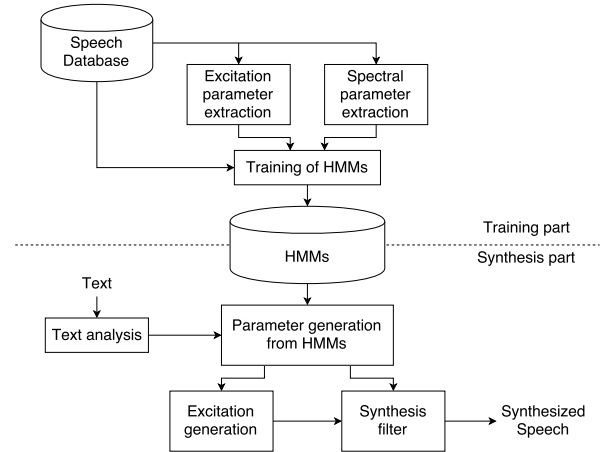

Figure 1: Function blocks of HMM-based synthesis [5]

In the training part spectrum and excitation parameters of the recorded speech are used to generate acoustic models represented by the HMMs. Therefore, phonetic, linguistic and prosodic parameters are considered. In comparison to a unit-selection system, the large speech database is only needed in the training part.

In the synthesis part, first the text which is to be synthesized is transformed to a sequence of parameters containing information about the context. According to this sequence, the respective HMMs are concatenated in order to form an utterance HMM. Then, after determining the state durations of the HMMs, a sequence of coefficients is created, which finally is used to construct the speech waveform using a specific filter (*e.g.* a Mel-Log Spectral Approximation (MLSA) filter).

The main disadvantage of this approach compared to unit-selection synthesis is the quality of the synthesized speech. Three factors are accountable for the lack of quality: the vocoder, the modeling accuracy, and an effect called over-smoothing.

However, there are some essential advantages, which make the HMM-based approach a competitive alternative to unit-selection

synthesis. First, the voice characteristics can be modified without much effort. Thus, the implementation of different languages and the realization of different speaking styles with emotional emphasis is possible. Second, these aspects require a much smaller database than in unit-selection synthesis, since only a statistic representation of speech segments rather than raw speech data is stored. Finally, the quality of the synthesized speech of a SPSS system is much more robust and free of sporadic flaws.

In Table 1, the techniques discussed so far are compared regarding the most prominent advantages and drawbacks.

Table 1: Comparison of speech generation methods [5, 12]

Technique	Advantages	Drawbacks
Formant-based	Very small footprint	Very artificial and metallic voice
Unit-selection	Very high voice quality possible	Large database required
HMM-based	Adjustable voice and small footprint	Voice sounds muffled

3 SPSS WITH DEEP LEARNING MODELS

Although the approach of SPSS has brought many advantages over unit-selection synthesis as shown in the second chapter, the generated voice is still not as natural as desired. Therefore, deep learning models recently have been used to further improve SPSS. Since Deep Neural Networks (DNNs) have proven to be very effective in speech recognition, they have found their way into speech synthesis, too [11].

In the beginning of this section, one specific approach how to introduce deep learning models into speech synthesis is explained further by using [17] as reference. Then, some alternative techniques of enhancing speech synthesis through the deployment of deep learning models are outlined, as shown in the work of Hashimoto *et al.* [11].

3.1 One Specific Approach for Improvement

In the previous section, I mentioned that the quality issue of HMM-based speech synthesis is caused by three aspects: the vocoder, the accuracy of acoustic models and the over-smoothing effect. Zen *et al.* [17] suggest a specific approach to eliminate one of these causes, *i.e.* the accuracy of acoustic models by allocating this task to a DNN. In conventional HMM-based systems the mapping between context features (phonetic and linguistic properties) and speech parameters is done by decision tree based context clustering. Thereby, the context-dependent HMMs are assigned to different clusters depending on the combination of contexts using binary decision trees. Each cluster is characterized by a specific set of speech parameters. In this way, it is possible to estimate all HMMs in a robust way with a typically sized training database.

However, decision trees soon reach their limits when handling complex contexts. Only by increasing the size and in this way decreasing the efficiency of the decision tree, more complicated contexts (*e.g.* XOR) can be dealt with. In addition to that, decision trees require partitioned input data with each partition having a different set of parameters. In this way, with less data per region overfitting is likely to happen, which then results in a lack of quality. These downsides can be avoided by using a DNN instead of multiple decision trees.

Nevertheless, the use of a DNN also introduces two disadvantages: The first one arises in terms of computational power. Both, in the training and in the prediction stage, decision trees require much less operations (total amount and level of complexity) than DNNs. The second one has to do with the decision process in its basic form.

With decision trees a binary question has to be answered, whilst a DNN consists of weighted neurons, which use non-linear activation functions (*e.g.* sigmoid, tanh, ReLU [8]) to determine their state. As consequence, interpretable rules are far easier to produce with decision trees than with DNNs.

In Figure 2, the structure of a DNN-based speech synthesis system is shown. First, a sequence of input features is generated after analyzing the input text. These parameters contain numeric values like the number of words in a sentence or the duration of a phoneme as well as binary answers to questions like "is the current phoneme *aa*?". Then, this parameter sequence is fed into the DNN where a mapping to output features is deployed by using forward propagation. The DNN has to be trained before with pairs of input and output features from a database. In the following steps, the speech parameters are extracted from the statistics of the output features and the voice waveform in turn generated from the speech parameters. This is done in the same way as in the HMM-based system. For this system the function blocks of text analysis, parameter generation, and waveform synthesis can be reused from an HMM-based system. Only the mapping from input features (*e.g.* linguistic contexts) to output features (spectral and excitation parameters) is implemented in a different way.
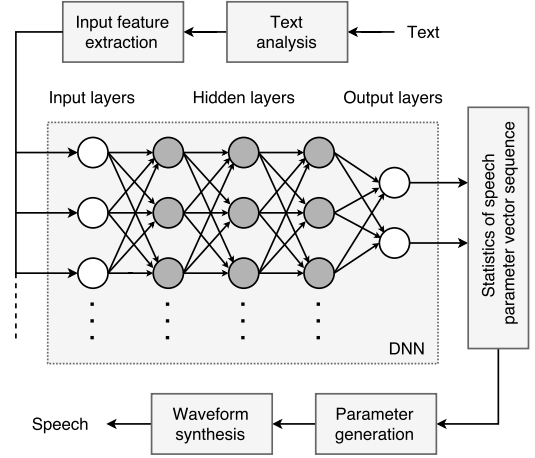

Figure 2: Speech synthesis based on DNN [17]

To compare the output of the above described framework with that of an HMM-based, Zen *et al.* conducted some experiments with each an HMM-based and a DNN-based speech synthesis system in [17]. Therefore, they used the same speech data in US English which includes about 33000 utterances for both systems. The HMM-based system used 2554 questions for the decision tree-based context clustering. To influence the size of the decision trees, the scaling factor α is used, where a high value results in a small decision tree and $\alpha = 1$ denotes a typical HMM-based system. 342 binary features (*e.g.* phonemes identities) and 25 numerical features (*e.g.* number of syllables in a word) represent the input features of the DNN-based system. The sigmoid function is used as activation function, since the authors experienced superior performance of this type in previous tests. In total, one network with different number of layers (1, 2, 3, 4, or 5) and units per layer (256, 1024, or 2048) are used.

The objective evaluation showed that the DNN-based system achieved better performance in voiced/unvoiced classification and aperiodicity prediction, regardless of the number of layers or the units per layer. At the Mel-cepstral distortion only DNNs with three or more layers outperformed the HMM-based system. As subjective evaluation, 173 test sentences each synthesized with the DNN- and the HMM-based systems were played back to a number of listeners,

who then choose which synthesis they preferred. If no difference is perceived, the option "neutral" can be selected. In this test, the same number of parameters for both systems are used. For the structure of the DNN-based approach, four layers with a different number of units per layer were applied. In Table 2, the outcome of the subjective evaluation is shown. Distinctly, the speech samples generated with the DNN-based systems were preferred, regardless of the amount of units per layer. The listeners described them as less muffled.

Table 2: Subjective scores (in %) of speech samples in [17]

HMM (α)	DNN (layers × units)	Neutral
15.8 (16)	38.5 (4 × 256)	45.7
16.1 (4)	27.2 (4 × 512)	56.8
12.7 (1)	36.6 (4 × 1024)	50.7

In conclusion, it can be stated that a DNN-based approach to implement the acoustic model of a speech synthesis system is a reasonable alternative to the conventional decision tree-based strategy. The improved performance for predicting spectral and excitation parameters and the more natural sounding voice both show the potential of the DNN-based approach for speech synthesis. However, the higher computational costs both at training and prediction stage due to more complex arithmetic operations in the DNN-based approach indicate where future work should be focused on.

3.2 Other Ways for Improvement

The previous section focused on the impact a DNN has, by representing the acoustic model of a speech synthesis system. However, deep learning models can be deployed in other parts of SPSS as well. In [11], different ways on how to implement a DNN into a speech synthesis are investigated. Therefor, different parts of an HMM-based speech synthesis system are once modeled with a DNN and once with the conventional technique. Then, the results are compared in an objective and a subjective fashion.

Hashimoto *et al.* concentrated on two core components of a speech synthesis system, the acoustic models and the speech generation part. In Figure 3, you can see the simplified structure of a speech synthesis system. In the first block, the text is analyzed and the contextual features are extracted (A). These are then converted to static and dynamic acoustic features (B) by the acoustic models. The parameter generation block then uses these acoustic features to create speech parameters (C). In the last step, the speech waveform is synthesized.


Figure 3: Simplified structure of speech synthesis system [11]

The two gray-colored blocks in Figure 3 are subject of two experiments conducted in [11]. The conventional approach for the acoustic models is the use of decision tree clustered HMMs, whereas the parameter generation usually is implemented by a Maximum Likelihood Parameter Generation (MLPG) algorithm. As seen in Section 3.1, the acoustic models can be represented by a DNN. Hashimoto *et al.* also included this approach in their experiments but furthermore used a DNN for the speech parameter generation task.

In Table 3, the systems resulting of the different combinations of the deployment of either conventional or DNN-based approach are shown. The systems I - IV are part of experiment 1 and all systems except system IV are part of experiment 2.

Table 3: Different systems within the experiments [11]

System Nr.	I	II	III	IV	V	VI
Experiments	1 & 2	1 & 2	1 & 2	1	2	2
Acoustic models	HMM	HMM	DNN	DNN	HMM	HMM
Speech generation	MLPG	DNN	MLPG	DNN	DNN + MLPG	MLPG + DNN

In experiment 1, all DNNs are equipped with three hidden layers and different number of units per layers (256, 512, or 1024), while the size of the decision trees was controlled similar as in [17]. Only the Mel-cepstral distortion was used as objective evaluation. Here, the systems with a DNN (system II - IV) all outperformed system I in a similar way, as soon as they had 512 or more units per layer. For the subjective evaluation, 20 sentences were played back to 8 listeners who had to assign a naturalness score between 1 and 5 (with 5 being the most natural). After applying the mean opinion score test method, system III was most preferred with 3.53 followed by system IV with 3.17 and system I with 3.08. The less natural sounding voice according to the listeners was produced by system II. From these results, it can be concluded that the speech generation part implemented with MLPG (system I & III) produces a more natural output than those with a DNN (system II & IV).

Due to these outcomes, experiment 2 introduces two new test cases, system V and system VI. For the parameter generation task a combination of MLPG and DNN is used, for which the respective input and output features were adapted. Again, once an objective and a subjective evaluation like in experiment 1 were conducted to measure the performance. The objective evaluation basically showed the same trend as in experiment 1, as soon as 512 or more units per layer were used, the Mel-cepstral distortion was smaller than with the conventional approach (system I). In the subjective evaluation, the two new systems (V & VI) both were preferred over system I, with system VI being the most preferred system of this experiment. System III again was better and system II again was evaluated less natural than system I.

As conclusion, we see that the experiments conducted in [11] confirmed the results of [17] that the decision tree-clustered HMMs can be replaced by a DNN to improve both the prediction performance as well as the naturalness of the generated speech. Beyond that Hashimoto *et al.* demonstrated that it is worthwhile to go one step further and deploy an additional DNN as post-filtering block into the parameter generation task of the speech synthesis system to further improve the above stated results.

4 SPEECH SYNTHESIS ON MOBILE DEVICES

In this section, two approaches to optimize speech synthesis for mobile devices are presented. In the first, some optimization steps on a conventional HMM-based system are suggested [15], whereas the second approach introduces deep learning models to get rid of the dependence on Internet-based services [6].

4.1 Motivation and Challenges

During the last 10 years, advances in technology have led to a tremendous spread of mobile devices, especially smartphones. According to [2], there are 2.1 billion smartphone users worldwide in 2016 which is forecasted to increase to almost 3 billion in 2020. In [2], it

also states that only in 2016 about 1.5 billion new smartphones were sold. From these numbers, we can derive that already today and with increasing significance in the near future, smartphones constitute an essential part of our daily life.

Especially on smartphones where the visual output is restricted to a rather small screen, usually not bigger than 5 inch in diagonal, speech enabled interaction can improve the user experience in a significant way. In [15], three application scenarios how speech interaction can assist are pointed out. Firstly, speech interaction can be used as an extension to existing communication channels. Secondly, if other outputs are restricted (*e.g.* while driving), speech can be deployed as main output. Thirdly, speech can help visually impaired or blind people to interact with a system at all (*e.g.* screen readers).

However, when implementing a speech synthesis application on a mobile device, the increased computational power and storage capacity on recent smartphones cannot be used as with desktop devices. One aspect is the power consumption. An application on a smartphone should use as less processing power as possible to avoid shortening the battery life unnecessarily [15]. Another issue is the restricted amount of main memory (16 - 32 MB RAM per app), which forces each application to be as resource-saving as possible. While accomplishing these demands, the output speech still needs to be processed in real-time to ensure a good user experience [6].

4.2 Optimized HMM-based Synthesis

One approach to optimize HMM-based speech synthesis for mobile devices can be seen in [15]. Here, Tóth *et al.* conduct a study to optimize an HMM-based speech synthesis system in terms of computational power, playback latency and footprint size while ensuring a reasonable quality of the synthesized speech. In the following, this approach is explained further.

As already pointed out in Section 2.1, HMM-based speech synthesis offers the best trade-off between speech quality and footprint size. This surely makes it the most appropriate approach to be deployed in mobile systems. Tóth *et al.* indicate that there already have been several approaches to optimize HMM-based speech synthesis. However, they claim to be the first who emphasize the optimization steps on resource constrained platforms to achieve a reduction in computational costs, while still fulfilling real-time responsiveness.

To measure the processing time of the speech synthesis procedure, Tóth *et al.* divide the TTS-process in three steps: The loading of the HMM database into main memory (1), the speech parameter generation (2), and the waveform synthesis (3). A 17 seconds sequence is chosen as test sentence, although it is unlikely that such a long utterance has to be synthesized often in real life (since in practice segments with one to five seconds are synthesized). That way, it can be ensured that shorter sequences are synthesized as desired.

The following four approaches to optimize HMM-based TTS are suggested by Tóth *et al.*:

- **A) Adjusting vocoder parameters**
- **B) Reducing the size of decision trees**
- **C) Introducing streaming synthesis**
- **D) Applying source code optimizations**

In the first approach, the conventional MLSA filter used for speech parameter generation is replaced by a Line Spectral Pair (LSP) filter. In this way, the performance in step (2) can be enhanced. Different orders of the LSP filter are tested to compare speech subjective quality, with an order of 18 being the conventional setup and 14, 12, and 10 being the optimization approaches.

The second approach dealt with the size of the decision trees. While lowering the number of nodes of a decision tree reduces the quality of synthesized speech, it also decreases the memory footprint

and the computational costs. Therefore, three different test settings were defined (see Table 4).

Table 4: Test settings of optimization approach B) [15]

	Nodes in decision trees	Footprint size
Original	6983	666 KB
Setting 1	4762	463 KB
Setting 2	2743	214 KB
Setting 3	1273	140 KB

General speech synthesis frameworks usually do not include the audio playback functionality to ensure platform independence. Implementing this part allows streaming synthesis. In this way, one already synthesized segment is played back while the next segment is generated. While long segments cause an undesired latency, too short segments result in uncontinuous audio playback. The optimal segment size is determined at run-time to achieve a compromise between low as possible latency and low as possible computational effort.

Low-resource devices differ from high-end devices not only in the amount of available memory and the computing performance, but also in the chip architecture. Memory management, fixed- versus floating-point calculation and conditional call management are points to be aware of. By implementing these issues in an architecture-optimized way, the processing load and time can be decreased.

The above described optimizations were tested on three different smartphones with increasing CPU speed and main memory, starting with the "weakest": an Apple Iphone 3G, a Samsung Galaxy Spica (GT-i5700) and a HTC Desire (A8181). Therefore, the above described optimization steps were tested one by one, since they do not affect each other in a noteworthy manner.

For the first approach a 12th order LSP filter was chosen as the optimal setting in terms of latency and subjective speech quality. Concerning the size of the decision trees, setting 1 clearly proved to be the optimal setting, since the time to load the HMM database (1) was decreased to half as before without notable impaired speech quality. The streaming synthesis approach tremendously decreased step (1), with an improvement from the order of seconds to 8 ms without affecting the speech quality at all. Concerning the source code optimizations, almost no optimizations have been conducted. Only some conditional calls were removed, which resulted in a insignificant improvement of latency. In total the computation time could be decreased by 65 % without a negligible decrease of speech quality.

4.3 Deep Learning-based Synthesis

The usual approach for current Virtual Personal Assistants (VPAs) (or generally Human Computer Interaction (HCI) interfaces) to use servers for processing input data and generating dedicated output actions, results in a total dependency on existing Internet access. In [6], Boros *et al.* are indicating three major disadvantages of this strategy: The availability of network access is not always granted (*e.g.* bad coverage, additional charges abroad), the issue of data security, and possible delays due to bad network connection causing bad user experience. The authors are therefore suggesting the use of deep learning models to port an existing TTS system (based on Romanian language) to a mobile device. Thereby, the main goal is to reduce the total memory footprint without loosing to much accuracy compared to the original system implementation.

A TTS system consists of a text processing front-end and a digital signal processing back-end. The suggestions made in [6] solely focus on the front-end. In this scope a deep learning model is introduced in each of the following text processing tasks:

1. **Syllabification (SYL)**
2. **Phonetic transcription (PT)**
3. **Part-of-speech Tagging (POT)**
4. **Lexical stress prediction (LSP)**

All experiments are conducted with a DNN with two or three hidden layers, where each hidden layer consists of an auto encoder. In the following, the experiments are explained further and the respective results are summarized in Table 5. As reference system, the Margin Infused Relaxed Algorithm (MIRA) is used [7].

In the syllabification task, each word is decomposed into syllables (phonological units). This contributes to the later conducted prosody generation. By using a data-driven approach instead of a rule-based approach, language independence is ensured. In comparison to the MIRA algorithm, the word-level accuracy of the DNN-based approach (two hidden layers) is only 0.78 % worse but enables a reduction of the model size by a factor of more than 250 to 36.7 KB.

Phonetic transcription is the process of translating letters to phonemes (also referred to as letter-to-sound or grapheme-to-phoneme) and is also chosen to be implemented in a data-driven approach, due to the same reason as in the syllabification. By using a DNN with two hidden layers, a word-level accuracy of 96.16 % can be achieved, which is a decrease of only 0.13 % compared to the MIRA-based technique. However, the model size again is reduced significantly from around 1.4 MB to 43.4 KB.

To determine how a word is spoken, its Part-of-Speech Tagging is one of the basic features. Allocating a uniquely interpretable tag to each word in a sentence is for example essential to choose how to pronounce a homograph[1]. In this part, two DNNs were deployed (both with two hidden layers), one to replace the lexicon containing the mapping of words to their syntactic information and one which performs the tagging instead of a standard feed-forward Neural Network (NN). In this task, the biggest reduction of the model size could be achieved, from almost 100 MB to about 178 KB. Thereby, the word-level accuracy got reduced by 3.03 %. This reduction is due to the lack of fine-tuning of the DNNs, which in this scope has not been conducted.

The last task is the lexical stress prediction, which represents an essential part of prosody generation, by indicating which syllables have to be given a special emphasis during the utterance. The DNN used for this task is equipped with three hidden layers and results in an overall accuracy of 97.67 % and a reduction of the model size by a factor of 60 to around 110 KB.

Table 5: Resulting accuracy and footprint size [6]

	SYL		PT		POT		LSP	
	MIRA	DNN	MIRA	DNN	NN	DNN	MIRA	DNN
Accuracy	99.01 %	98.23 %	96.29 %	96.16 %	98.19 %	95.16 %	98.80 %	97.67 %
Size	9.4 MB	36.7 KB	1.4 MB	43.4 KB	96 MB	178 KB	6 MB	110 KB

Overall, it can be concluded that the measures introduced in [6] enabled a tremendous reduction of the model sizes, while keeping the performance almost as high as in the previous models. From these results, it can be deduced that the introduction of deep learning models like DNNs into the front-end or text processing part of a TTS-system is a highly reasonable strategy to port speech synthesis systems onto mobile devices without relying anymore on Internet access.

5 CONCLUSIONS

In this paper, a systematic review of the impact of deep learning models on speech synthesis on mobile devices is given. Therefore,

different approaches to improve the conventional HMM-based synthesis are pointed out, including the deployment of DNNs as acoustic models or in the parameter generation task. Here, a significant improvement of the prediction performance and the speech quality could be achieved [11, 17]. In the following, the goal of modifying speech synthesis for the use in resource-constrained environments, such as mobile devices, is outlined by reference to two strategies: one with and one without the use of deep learning models. In the first one, the focus is set on adjusting different parameters and applying several optimization steps [15] in order to achieve real-time playback. As a result, the computation time could be decreased by 65 % without a negligible loss of speech quality. The last approach [6] tackled the dependency of network access, while using speech synthesis applications on mobile devices. Therefore the authors suggested the use of DNNs in several parts of the front-end of a TTS system so as to reduce the footprint size, which enables the offline use of speech synthesis applications.

It is expected that the acceptance of VPAs like Apple's Siri or Amazon's Alexa will increase in the near future [4]. Therefore, the development of robust and resource-efficient speech synthesis methods is an essential part to meet the challenges of mobile environments. Innovative technologies like personalized speech-to-speech translation or voice cloning are only two examples of emerging techniques, wherefore speech synthesis has to be as evolved as possible [16].

REFERENCES

[1] 2017. The Benefits of Text to Speech. (02 June 2017). http://www.readspeaker.com/benefits-of-text-to-speech/

[2] 2017. Number of smartphone users worldwide. (10 June 2017). https://www.statista.com/statistics/330695/number-of-smartphone-users-worldwide/

[3] 2017. Stephen Hawking - The Computer. (02 June 2017). http://www.hawking.org.uk/the-computer.html

[4] 2017. Survey on usage of virtual personal assistants by Gartner. (22 May 2017). http://www.gartner.com/newsroom/id/3551217

[5] A. W. Black, H. Zen, and K. Tokuda. 2007. Statistical Parametric Speech Synthesis. In *2007 IEEE International Conference on Acoustics, Speech and Signal Processing - ICASSP '07*, Vol. 4. IV–1229–IV–1232. https://doi.org/10.1109/ICASSP.2007.367298

[6] T. Boroş and S. D. Dumitrescu. 2015. Robust Deep-learning Models for Text-to-speech Synthesis Support on Embedded Devices. In *Proceedings of the 7th International Conference on Management of Computational and Collective intElligence in Digital EcoSystems (MEDES '15)*. ACM, New York, NY, USA, 98–102. https://doi.org/10.1145/2857218.2857234

[7] T. Boros. 2013. A unified lexical processing framework based on the Margin Infused Relaxed Algorithm. A case study on the Romanian Language. In *Recent Advances in Natural Language Processing, RANLP 2013, 9-11 September, 2013, Hissar, Bulgaria*. 91–97. http://aclweb.org/anthology/R/R13/R13-1012.pdf

[8] H. Chung, S. J. Lee, and J. G. Park. 2016. Deep neural network using trainable activation functions. In *2016 International Joint Conference on Neural Networks (IJCNN)*. 348–352. https://doi.org/10.1109/IJCNN.2016.7727219

[9] L. Deng. 2014. A tutorial survey of architectures, algorithms, and applications for deep learning. *APSIPA Transactions on Signal and Information Processing* 3 (January 2014). https://doi.org/10.1017/atsip.2013.9

[10] L. Dhavala. 2014. Communication between Air Traffic Controllers and Pilots. (November 2014). https://doi.org/10.13140/2.1.3977.8887

[11] K. Hashimoto, K. Oura, Y. Nankaku, and K. Tokuda. 2015. The effect of neural networks in statistical parametric speech synthesis. In *2015 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 4455–4459. https://doi.org/10.1109/ICASSP.2015.7178813

[12] F. Hinterleitner. 2017. *Quality of Synthetic Speech: Perceptual Dimensions, Influencing Factors, and Instrumental Assessment (T-Labs Series in Telecommunication Services)* (1st ed. 2017 ed.). Springer, Singapore. https://doi.org/10.1007/978-981-10-3734-4

[13] Z. H. Ling, S. Y. Kang, H. Zen, A. Senior, M. Schuster, X. J. Qian, H. M. Meng, and L. Deng. 2015. Deep Learning for Acoustic Modeling in Parametric Speech Generation: A systematic review of existing techniques and future trends. *IEEE Signal Processing Magazine* 32, 3 (May 2015), 35–52. https://doi.org/10.1109/MSP.2014.2359987

[14] D. Suendermann, H. Höge, and A. Black. 2010. *Challenges in Speech Synthesis*. Springer US, Boston, MA, 19–32. https://doi.org/10.1007/978-0-387-73819-2_2

[15] B. Tóth and G. Németh. 2012. Optimizing HMM Speech Synthesis for Low-Resource Devices. In *Journal of Advanced Computational Intelligence and Intelligent Informatics*. 327–334.

[16] J. Yamagishi. 2011. New and emerging applications of speech synthesis. (February 2011). www.cstr.ed.ac.uk

[17] H. Zen, A. Senior, and M. Schuster. 2013. Statistical parametric speech synthesis using deep neural networks. In *2013 IEEE International Conference on Acoustics, Speech and Signal Processing*. 7962–7966. https://doi.org/10.1109/ICASSP.2013.6639215

---

[1]A word with different meanings depending on the pronunciation